Artificial Intelligence — Introduction

Quick offline notes (study-ready formulae + key ideas):

Loss measures prediction error
Gradient descent minimizes loss over parameters
Overfitting: training good, validation poor
Regularization: reduce complexity, improve generalization
Classification: choose class with highest score
Don't forget: train/val/test split

Tip: open the PDF and try 3 practice questions from your course notes.